


Pontificia Universidad Católica de Chile  
Escuela de Ingeniería  
Departamento de Ciencia de la Computación



## IIC1005 – Computación: Ciencia y Tecnología del Mundo Digital

### Introduction to Deep Learning



# ¿Why a class solely focused on Deep Learning?

(aka why is Deep Learning so important?)


AlphaGo seals 4-1 victory over world Go champion Lee Sedol

DeepMind's artificial intelligence astonishes as it offers evidence computer software has



## Rise of the machines: Google AI experiment may lead to robots that can learn WITHOUT human input


- Generative Adversarial Networks create digital content based on real-life
- Google project pits AI algorithms against each other to refine this output
- The results could one day lead to machines that can learn without human input



Not new course to

now and being trialled in south-east London now  
be banned from driving?





## Artificial Intelligence Takes Off at Google

Number of software projects within Google that uses a key AI technology, called Deep Learning.



Source: Google

Note: 2015 data does not incorporate data from Q4



**Andrej Karpathy** ✓

@karpathy

Following

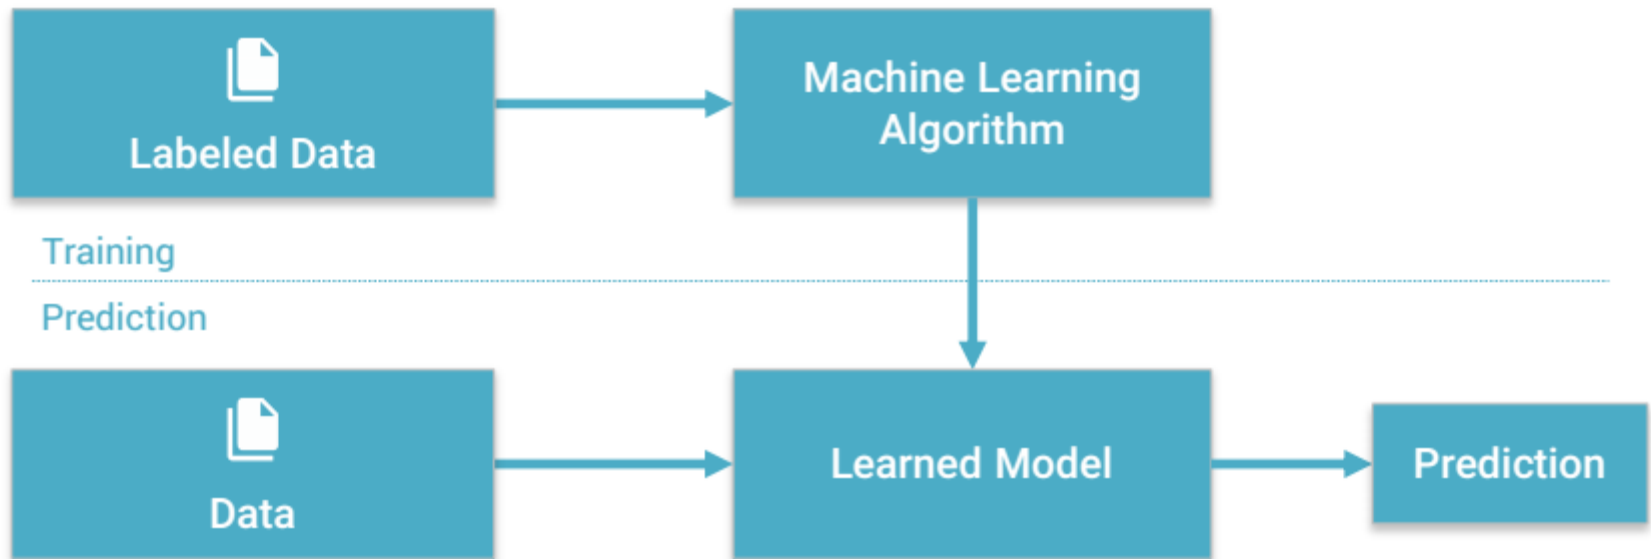


Came to visit first class of [@cs231n](#) at Stanford. 2015: 150 students, 2016: 350, this year: 750. [#aiinterestsingularity](#)





**Machine Learning** is a type of Artificial Intelligence that provides computers with the ability to **learn without being explicitly programmed**.



Provides **various techniques** that can learn from and make predictions on data

# Deep Learning is currently the most popular tool/technique in Machine Learning



Part of the machine learning field of learning representations of data. Exceptional effective at learning patterns.



Utilizes learning algorithms that derive meaning out of data by using a hierarchy of multiple layers that mimic the neural networks of our brain.



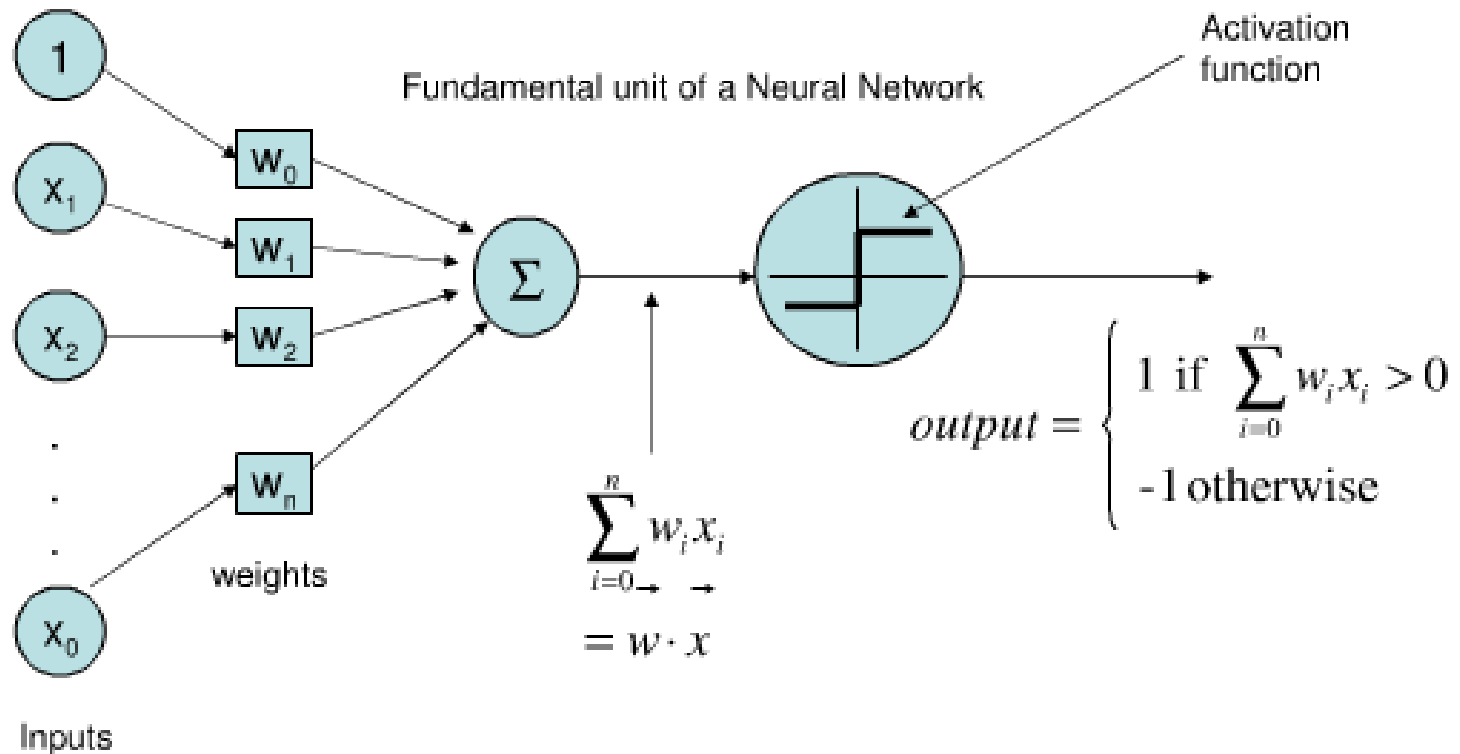
If you provide the system tons of information, it begins to understand it and respond in useful ways.

Deep Learning's power comes from its ability to learn a hierarchy of meaningful representations from data







## All started with a tiny Perceptron



...but it wasn't very powerful



Linearly separable



Non-Linearly separable



**1958** Perceptron

**1974** Backpropagation

Convolution Neural Networks for  
Handwritten Recognition



**1998**

Google Brain Project on  
16k Cores



**2012**

awkward silence (AI Winter)


**1969**

Perceptron criticized



**1995**

SVM reigns



**2006**


Restricted  
Boltzmann  
Machine




**2012**

AlexNet wins  
ImageNet

IMAGENET



A deep neural network consists of a **hierarchy of layers**, whereby each layer **transforms the input data** into more abstract representations (e.g. edge  $\rightarrow$  nose  $\rightarrow$  face). The output layer combines those features to make predictions.



## Unbeatable in many domains



Speech  
Recognition



Computer  
Vision




Natural Language  
Processing

## The right elements at the right time



Big Data  
(Digitalization)



Computation  
(Moore's Law, GPUs)



Algorithmic  
Progress

There won't be a new "AI winter" (hopefully)

facebook



YAHOO!

Google



IBM




NVIDIA®

Baidu 百度




## Let's get serious



Consists of one input, one output and multiple fully-connected hidden layers in-between. Each layer is represented as a series of neurons and progressively extracts higher and higher-level features of the input until the final layer essentially makes a decision about what the input shows. The more layers the network has, the higher-level features it will learn.

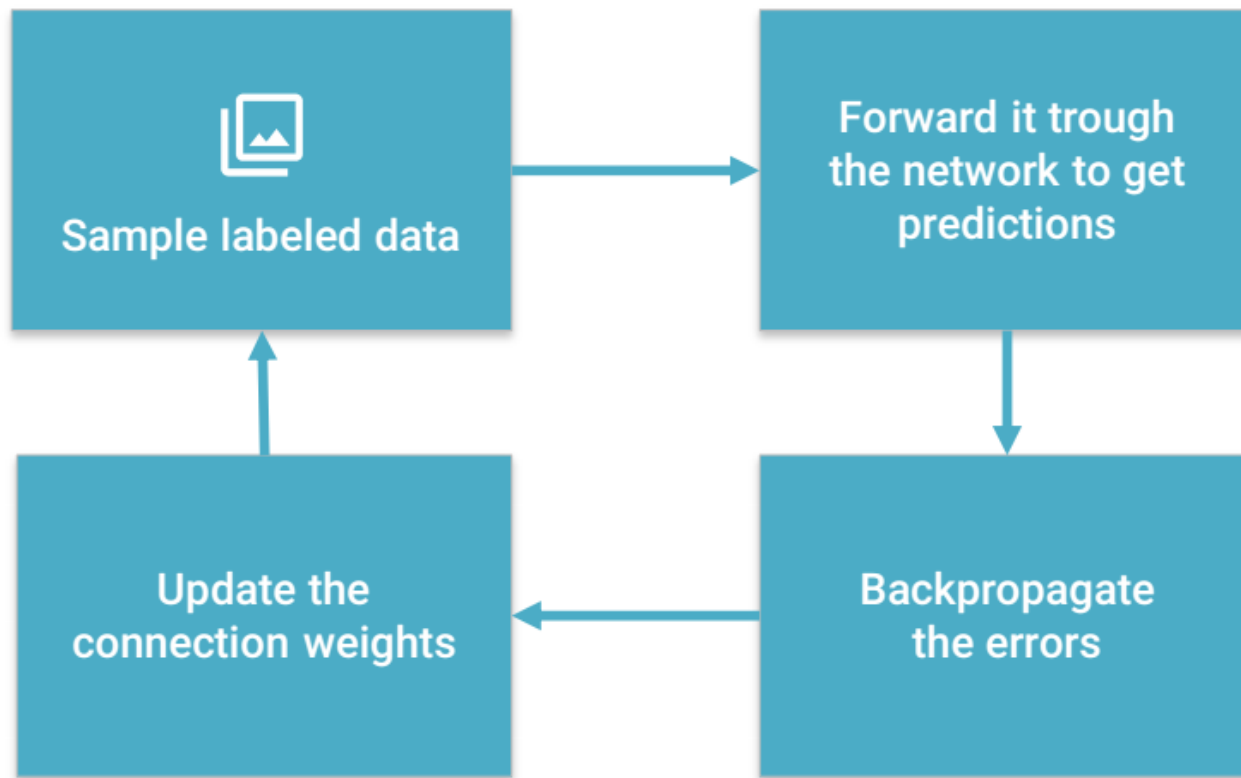
# Artificial neurons are somewhat similar to real neurons



An artificial neuron contains a **nonlinear activation function** and has several incoming and outgoing **weighted connections**.




Neurons are **trained to filter and detect specific features** or patterns (e.g. edge, nose) by receiving weighted input, transforming it with the activation function and passing it to the outgoing connections.



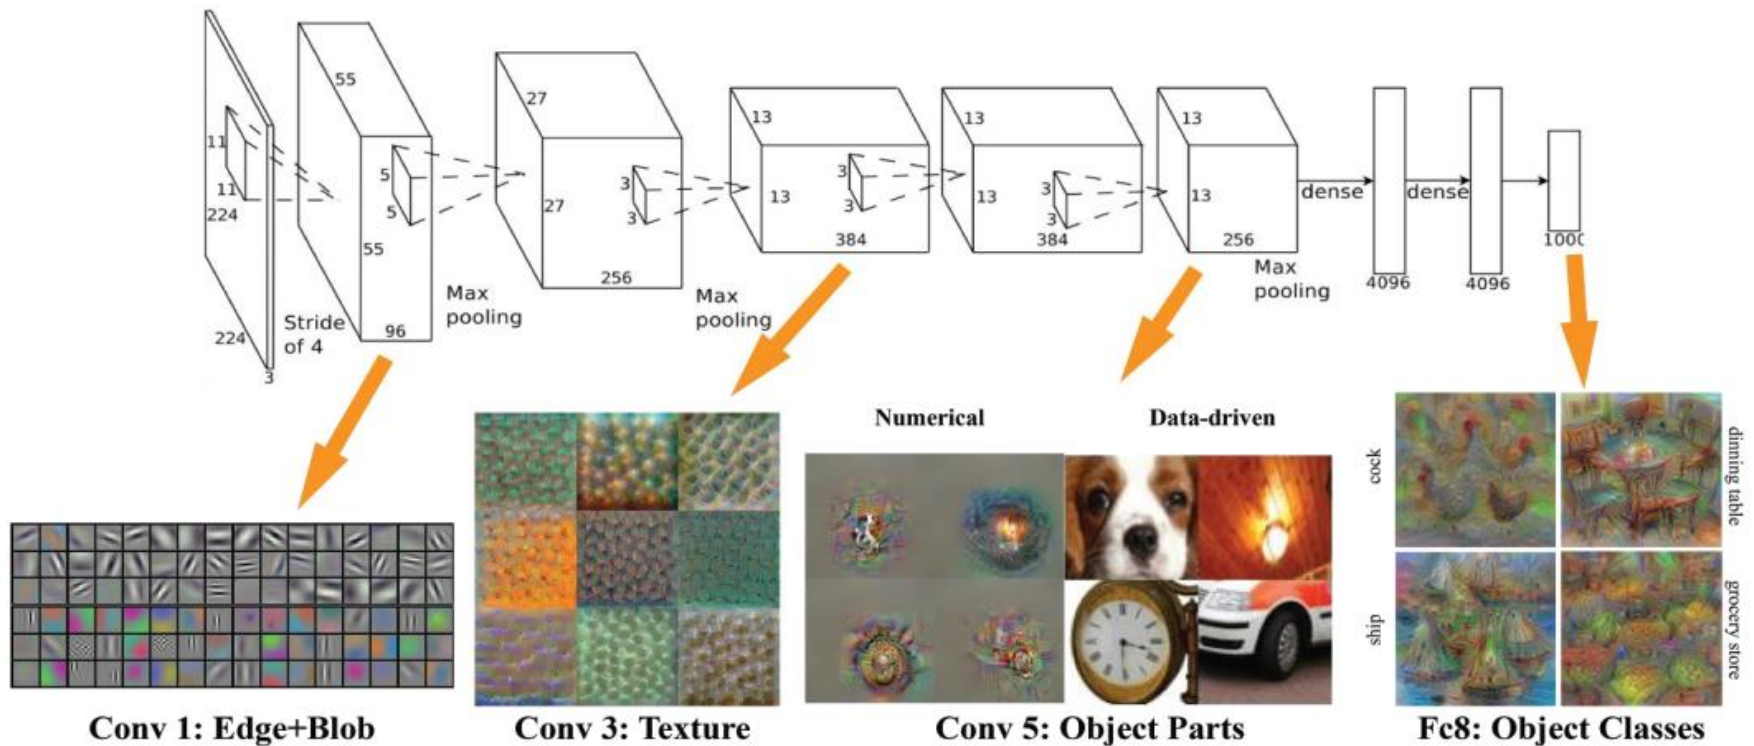
Learns by generating an error signal that measures the difference between the predictions of the network and the desired values and then **using this error signal to change the weights** (or parameters) so that predictions get more accurate.

Convolutional Neural Networks learn a complex representation of visual data using vast amounts of data. They are **inspired by the human visual system** and learn **multiple layers of transformations**, which are applied on top of each other to extract a progressively more **sophisticated representation of the input**.



Every layer of a CNN **takes a 3D volume of numbers and outputs a 3D volume of numbers**. E.g. Image is a 224\*224\*3 (RGB) cube and will be transformed to 1\*1000 vector of probabilities.






Convolution layer is a feature detector that automatically learns to **filter out not needed information** from an input by using convolution kernel.


Pooling layers compute the max or **average value of a particular feature over a region** of the input data (*downsizing of input images*). Also helps to detect objects in some unusual places and reduces memory size.

## If there is no supervision available, we can use Autoencoders




Composed of two symmetrical deep-belief networks. The encoding network learns to compresses the input to a condensed vector (dimensionality reduction). The decoding network can be used to reconstruct the data.


one to one




one to many




many to one



many to many



many to many




RNNs are **general computers which can learn algorithms to map input sequences to output sequences** (flexible-sized vectors). The output vector's contents are influenced by the entire history of inputs.



**State-of-the-art results** in time series prediction, adaptive robotics, handwriting recognition, image classification, speech recognition, stock market prediction, and other sequence learning problems.  
*Everything can be processed sequentially.*






To train the RNN, insert characters sequentially and predict the probabilities of the next letter. Backpropagate error and update RNN's weights to increase the confidence of the correct letter (green) and decrease the confidence of all other letters (red).

The emperor travelled back to [[Antioch, Perth, October 25|21]] to note, the Kingdom of Costa Rica, unsuccessful fashioned the [[Thrales]], [[Cynth's Dajoard]], known in western [[Scotland]], near Italy to the conquest of India with the conflict. Copyright was the succession of independence in the slop of Syrian influence that was a famous German movement based on a more popular servicious, non-doctrinal and sexual power post. Many governments recognize the military housing of the [[Civil Liberalization and Infantry Resolution 265 National Party in Hungary]], that is sympathetic to be to the [[Punjab Resolution]]  
(PJS)[<http://www.humah.yahoo.com/guardian.cfm/7754800786d17551963s89.htm>]


Trained on structured Wikipedia markdown. Network learns to spell English words completely from scratch and copy general syntactic structures.



FRANCE	JESUS	XBOX	REDDISH	SCRATCHED	MEGABITS
AUSTRIA	GOD	AMIGA	GREENISH	NAILED	OCTETS
BELGIUM	SATI	PLAYSTATION	BLUISH	SMASHED	MB/S
GERMANY	CHRIST	MSX	PINKISH	PUNCHED	BIT/S
ITALY	SATAN	IPOD	PURPLISH	POPPED	BAUD
GREECE	KALI	SEGA	BROWNISH	CRIMPED	CARATS
SWEDEN	INDRA	PSNUMBER	GREYISH	SCRAPED	KBIT/S
NORWAY	VISHNU	HD	GRAYISH	SCREWED	MEGAHERTZ
EUROPE	ANANDA	DREAMCAST	WHITISH	SECTIONED	MEGAPIXELS
HUNGARY	PARVATI	GEFORCE	SILVERY	SLASHED	GBIT/S
SWITZERLAND	GRACE	CAPCOM	YELLOWISH	RIPPED	AMPERES




Woman – Man  $\approx$  Aunt - *Uncle*  
 King - Male + Female  $\approx$  *Queen*  
 Human - Animal  $\approx$  *Ethics*



**A group of people shopping at an outdoor market.**

**There are many vegetables at the fruit stand.**

Neural Image Caption Generator **generates fitting natural-language captions only based on the pixels** by combining a vision CNN and a language-generating RNN.



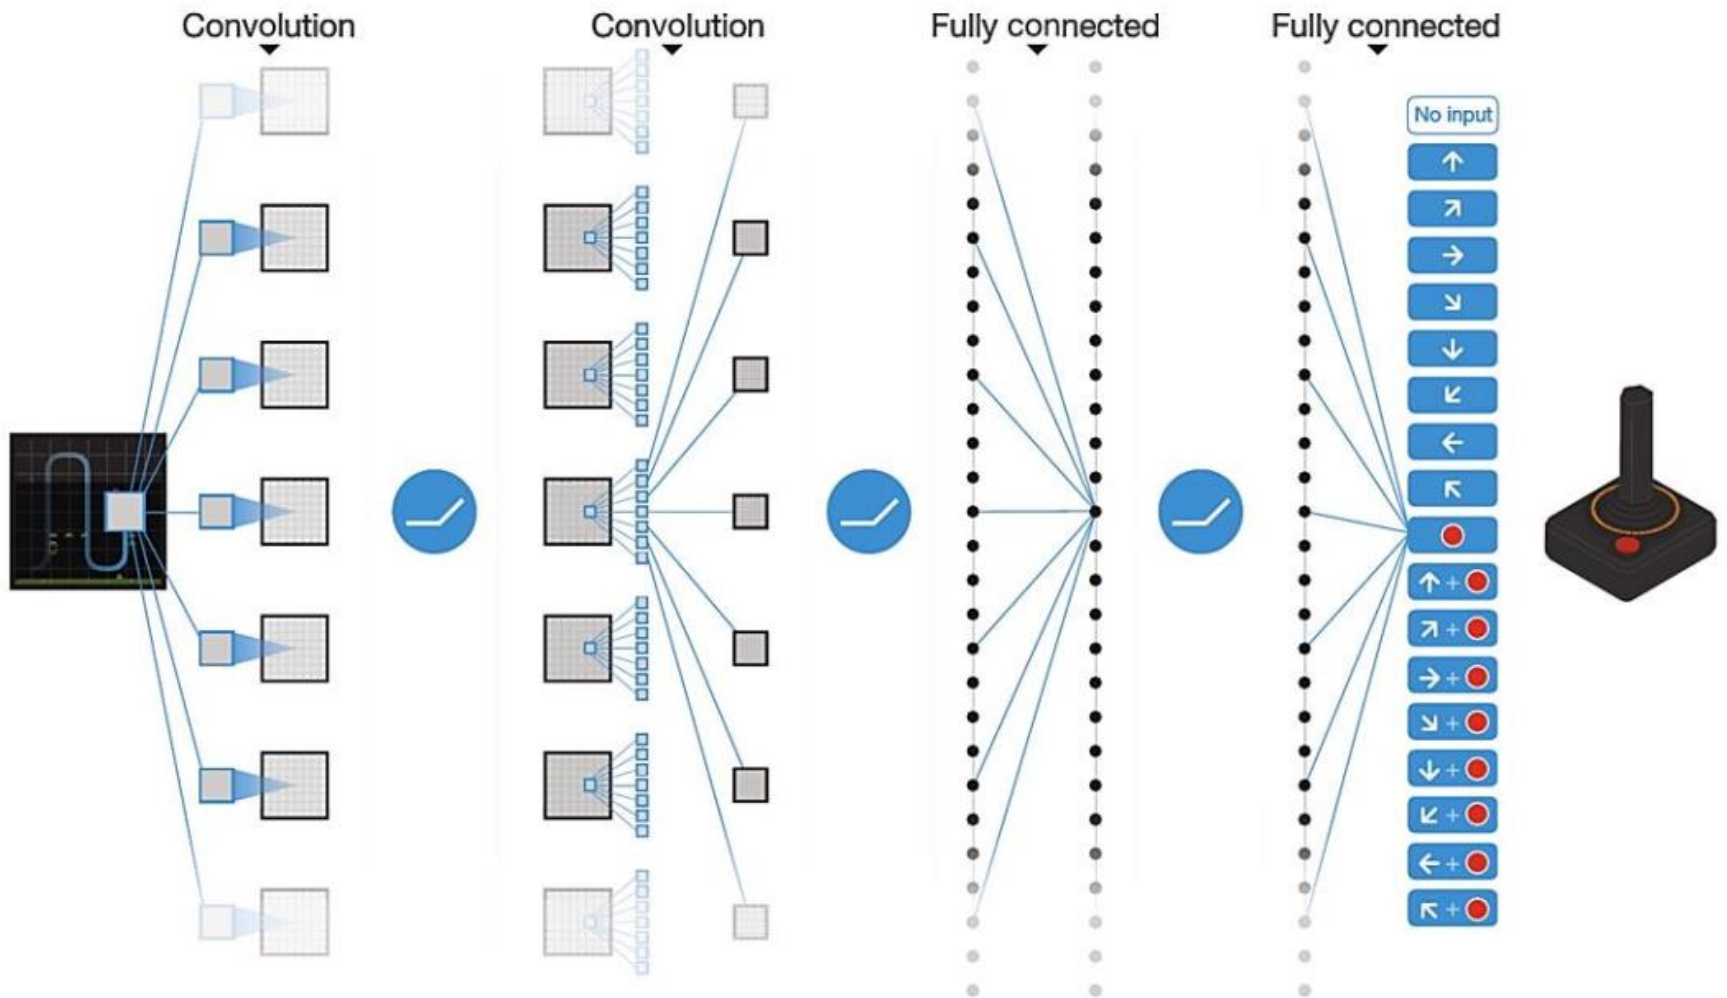
A close up of a child holding a stuffed animal



Two pizzas sitting on top of a stove top oven




A man flying through the air while riding a skateboard



Deep Q-Learning (DQN) is a model-free approach to reinforcement learning using deep networks in environments with discrete action choices


Ok, I'm sold, what should I do next  
(besides taking Machine Learning courses)



Read about it  
(serious reading)



MOOCs



## Most frameworks are open source

Caffe



DL4J  
Deeplearning4j

MINERVA

*mxnet*



MatConvNet



theano





# Great team, great professors, great research



Domingo Mery



Álvaro Soto



Jorge Baier



Karim Pichara



Cristian Ruz



Denis Parra



Hans Löbel

Great team, great professors, great research



# Takeaways

- Deep Learning is not magic! Just statistics and optimization in a black box.
- Exceptionally effective at recognizing patterns.
- Still no common sense (can be even racist).
- If you're interested, learn it now.
- If not, learn it anyways, at least the basics.
- This is one of the technologies that will change the world.



Pontificia Universidad Católica de Chile  
Escuela de Ingeniería  
Departamento de Ciencia de la Computación



## IIC1005 – Computación: Ciencia y Tecnología del Mundo Digital

### Introduction to Deep Learning

